

# LinguaFusion: Multi-Source Linguistic Fusion for News Classification

Pushing Accuracy Beyond TF-IDF with Interpretable Linguistic Features

---

Duo Zhang, Shuyuan Yang, Veronica Zhao

May 5, 2025

New York University

1. Motivation: Beyond Deep Learning Complexity
2. Experimental Setup
3. Linguistic Feature Exploration
4. Distributional Semantics
5. Cumulative Fusion Results
6. Conclusion

# Overview: Multi-Source Linguistic Fusion


## Key Question:

- How far can we improve document classification with simple, interpretable models?

## Our Approach:

- Combine five complementary linguistic signals
- Use simple logistic regression classifier
- Maintain full interpretability

**Result:** +3.32% accuracy improvement over TF-IDF baseline



## **Motivation: Beyond Deep Learning Complexity**

---

# The Challenge with State-of-the-Art

## Deep models (BERT, Transformers):


- **High Computational Cost:** Significant resources for training and inference.
- **Complex Architecture:** Many layers, millions of parameters.
- **Lack of Interpretability:** "Black boxes" make error analysis difficult.

## Our Goal

Push accuracy beyond a strong TF-IDF baseline using shallow, computationally cheap, and interpretable models.

## Research Question

How far can we improve accuracy by fusing complementary linguistic signals (POS, NER, GloVe, Doc2Vec) with a simple Logistic Regression classifier?



# Experimental Setup

---

# Data and Evaluation Protocol

## Corpus:


- 20 Newsgroups Dataset
- $\approx 18,000$  posts
- 20 balanced topics

## Baseline Model:

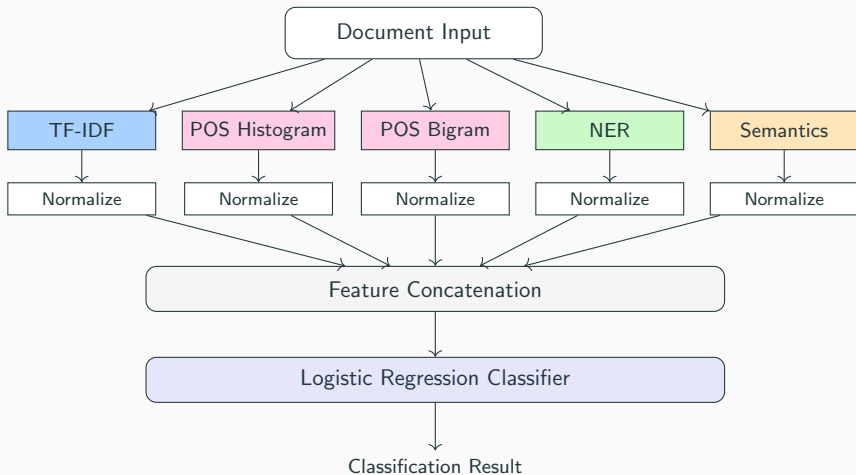
- TF-IDF + Logistic Regression
- Hyperparameters tuned on dev set, then fixed for all experiments

## Evaluation:

- 60% Train / 40% Test split
- Within 60 % train, 80 % for fitting / 20 % for development
- Final testing on hold-out set
- Metrics:
  - Macro-Averaged Accuracy
  - Per-class P/R/F1



# LinguaFusion Pipeline





# Linguistic Feature Exploration

---

# POS Features: Capturing Syntactic Fingerprints

## Two approaches to leverage syntax:


- **POS Histogram:** Normalized counts of 45 Penn Treebank tags. A 45-dim "syntactic fingerprint".
- **POS Bigram TF-IDF:** TF-IDF applied to sequences of two POS tags (e.g., NNP→VBZ). Captures local syntactic structure ( $\approx$  1500 dims after SVD).

Example:

*"The president announced policy."*

DT → NNP → VBD → NN

Tag Distribution:



## Impact


- Provide targeted lifts in classes with distinctive syntax (e.g., sports, politics).
- Histograms capture global tag mix; Bigrams capture local style.
- POS histogram boosted precision in 8 out of 20 categories, and POS bigrams lifted precision in 11 out of 20.

# POS Patterns: Examples from 20 Newsgroups

Topic	Distinctive POS Patterns
rec.sport.baseball	High VB, VBD (action verbs), NNP (team/player names) <i>"Martinez <b>hit</b> the ball over the fence"</i>
sci.electronics	High NN, JJ (technical terms with adjectives) <i>"The <b>digital circuit</b> requires precise voltage"</i>
talk.politics.mideast	Complex NP, nested clauses (formal discourse) <i>"The <b>committee</b> on <b>foreign affairs</b> has decided..."</i>

## Key Insight

POS histograms reveal that sports discussions have more action verbs (VBD), tech forums have more nouns (NN), and political discussions have more determiners (DT) and prepositions (IN).



# Illustrative Python Code: pos\_histogram

Function to compute normalized POS tag counts:

```
SELECTED_TAGS = ["CC", "CD", "DT", "EX", "FW", "IN", "JJ", "JJR", "JJS", "LS", "MD", "NN", "NNS", "NNP",  
"NNPS", "PDT", "POS", "PRP", "PRP$", "RB", "RBR", "RBS", "RP", "SYM", "TO", "UH", "VB", "VBD", "VBG", "VBN",  
"VBP", "VBZ", "WDT", "WP", "WP$", "WRB", "#", "$", "'", ",", ".", ":", "-LRB-", "-RRB-"]
```

```
TAG2IDX = {t:i for i,t in enumerate(SELECTED_TAGS)}
```

```
def pos_histogram(doc: str):  
    vec = np.zeros(len(SELECTED_TAGS), dtype=np.float32)  
    for _, tag in nltk.pos_tag(nltk.word_tokenize(doc)):  
        idx = TAG2IDX.get(tag)  
        if idx is not None:  
            vec[idx] += 1  
    if vec.sum():  
        vec /= vec.sum()  
    return vec
```

```
pos_hist_block = Pipeline([  
    ("hist", FunctionTransformer(lambda docs: np.vstack([pos_histogram(d) for d in docs]), validate=False)),  
    ("scal", StandardScaler()),  
)
```

# NER Features: Identifying Key Entities

## Entity Type Distribution:


- 9-dimensional histogram of entity types (PERSON, ORG, GPE, DATE, etc.)
- Normalized to create an "entity fingerprint"
- Added to TF-IDF vectors

**Example:** "Apple released the new iPhone in California."

Entity types: [ORG, PRODUCT, GPE]

## Impact

- NER features provide targeted lifts in classes with distinctive entity patterns (e.g., organizations in tech, people in sports).
- NER feature boosted precision in 8 out of 20 categories.



# Named Entity Histogram

## Method

- Count occurrences of 9 most frequent entity types per document
- Normalize counts to get a 9-dim NER histogram
- Concatenate to TF-IDF vector

```
NER_TYPES = ["PERSON","NORP","FAC","ORG","GPE","LOC", "PRODUCT","EVENT","DATE","TIME","CARDINAL","MONEY"]
NER2IDX = lbl:i for i,lbl in enumerate(NER_TYPES)
class NerHistogram(BaseEstimator,TransformerMixin):
    def fit(self,X,y=None): return self
    def transform(self,X):
        out=np.zeros((len(X),len(NER_TYPES)),dtype=np.float32)
        for i, doc in enumerate(_NLP.pipe(X, batch_size=32)):
            for ent in doc.ents:
                j = NER2IDX.get(ent.label_)
                if j is not None: out[i, j] += 1
        row_sums = out.sum(axis=1)
        nonzero = row_sums > 0
        out[nonzero] = out[nonzero] / row_sums[nonzero, np.newaxis]
        return out
ner_block = Pipeline([("hist",NerHistogram()), ("scal",StandardScaler())])
```


# Distributional Semantics

---

# Distributional Semantics: Capturing Meaning


## GloVe Embeddings:

- Pre-trained on large corpora
- Word-level vectors (300-dim) based on global co-occurrence
- Captures semantic relationships
- *Example:* king - man + woman  $\approx$  queen
- *Method:* Average vectors of words in document



## Doc2Vec Embeddings:

- Trained directly on 20 Newsgroups corpus
- Document-level vectors ("paragraph vectors")
- Learns representations capturing topic/style
- *Method:* Use Gensim's implementation



## Impact (Hybrid: TF-IDF + GloVe + Doc2Vec):

+2.87 pp test accuracy vs TF-IDF baseline, precision gain in 15 out of 20 categories.



## Cumulative Fusion Results

---

# Putting It All Together


Question: Do POS and NER features still add value to our distributional-semantics-enriched hybrid model (TF-IDF + GloVe + Doc2Vec)?

## Method: Simple Concatenation

Concatenate normalized/standardized feature blocks:

- TF-IDF (surface text)
- GloVe (averaged word embeddings)
- Doc2Vec (document embeddings)
- POS Histogram (45-dim)
- POS Bigram TF-IDF ( $\approx 1.5\text{k-dim}$ )
- NER Histogram (9-dim)

Feed the combined vector into the **same fixed Logistic Regression classifier**.



# Feature Union Code (Simplified)

```
tfidf_block = Pipeline([
    ("tfidf", TfidfVectorizer(stop_words="english",
                             max_features=best_tfidf_params["tfidf__max_features"],
                             ngram_range=best_tfidf_params["tfidf__ngram_range"])),
    ("scal", StandardScaler(with_mean=False))
])

pos_hist_block = Pipeline([...])
pos_bi_block = Pipeline([...])
ner_block = Pipeline([...])
glove_block = Pipeline([...])
doc2vec_block = Pipeline([...])

hybrid_model = Pipeline([
    ("features", FeatureUnion([
        ("tfidf", tfidf_block),
        ("glove", glove_block),
        ("d2v", doc2vec_block),
        ("pos_hist", pos_hist_block),
        ("pos_bi", pos_bi_block),
        ("ner", ner_block)
    ])),
    ("clf", LogisticRegression(max_iter=1000, C=1))
])
```


# Comparison of Feature Combinations (Test Set)

Model Configuration	Dev Acc.	Test Acc.	vs TF-IDF Baseline	
			Abs. $\Delta$	Rel. $\Delta$
TF-IDF (Baseline)	0.8944	0.8216	—	—
TF-IDF + POS (Hist)	0.8992	0.8139	-0.0077	-0.94%
TF-IDF + POS (Bigram)	0.9037	0.8184	-0.0032	-0.39%
TF-IDF + POS (Both)	0.9054	0.8186	-0.0030	-0.37%
TF-IDF + NER	0.9001	0.8103	-0.0113	-1.38%
TF-IDF + POS (Both) + NER	0.9041	0.8200	-0.0016	-0.19%
GloVe (Standalone)	0.7698	0.7098	-0.1118	-13.61%
Doc2Vec (Standalone)	0.8484	0.7675	-0.0541	-6.58%
Hybrid (TF-IDF+GloVe+D2V)	0.9205	0.8448	+0.0232	+2.82%
Hybrid + POS (Hist)	0.9218	0.8460	+0.0244	+2.97%
Hybrid + POS (Bigram)	0.9196	<b>0.8489</b>	<b>+0.0273</b>	<b>+3.32%</b>
Hybrid + POS (Both)	0.9200	0.8486	+0.0270	+3.29%
Hybrid + NER	0.9183	0.8456	+0.0240	+2.92%
Hybrid + POS (Both) + NER	0.9187	0.8481	+0.0265	+3.23%

## Key Finding

- The Hybrid model significantly outperforms TF-IDF.
- Adding POS/NER features to the Hybrid model provides further incremental gains.
- POS Bigrams provide the highest overall improvement (+3.32%).


# Results Visualization



# Why Syntax & Entities Still Matter

Even with powerful distributional embeddings (GloVe, Doc2Vec), POS and NER features add value:

- **Non-redundant Signals:** Embeddings capture *what* words mean; POS/NER capture *how* they are used (syntax, structure, salience).
- **Error Diversity / Orthogonal Cues:** Help disambiguate semantically similar words used in different contexts (e.g., "game" in sports vs. video games).
- **Lightweight:** Computationally cheap to extract and add minimal dimensionality instead of stacking deep architectures.



Example: "game"  
"The game was exciting"  
↓  
Sports context  
  
"This game has graphics"  
↓  
Computing context

## Conclusion

---


# Conclusion: LinguaFusion Takeaways

## Main Finding

Strategic fusion of lightweight, interpretable linguistic features (TF-IDF, POS, NER, GloVe, Doc2Vec) significantly boosts classification accuracy (+3.32 pp) over a strong TF-IDF baseline, without resorting to complex deep learning models.

## Key Insights:

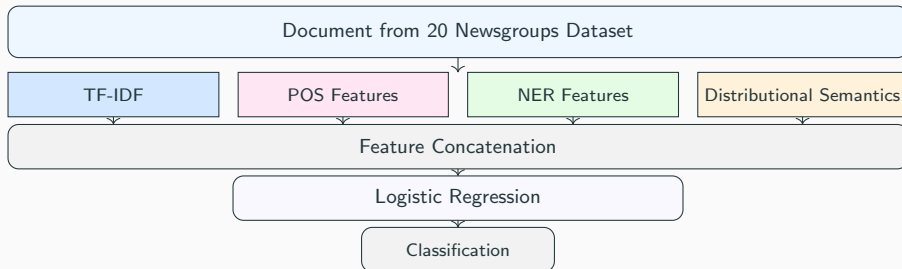
- **Complementarity is Key:** Different feature types capture distinct linguistic aspects; their combination is powerful.
- **Lightweight Block:** Uses inexpensive feature blocks, enabling fast training and serving on modest hardware.
- **Interpretability is Preserved:** Each feature block has a clear linguistic meaning, aiding analysis and debugging.





# LinguaFusion: Architecture Detail

Feature Block	Extractor	Dimensions	Library
TF-IDF	TfidfVectorizer	$\approx 20,000$	sklearn
POS Histogram	POS tagger	45	NLTK
POS Bigrams	POS tagger + TF-IDF	$\approx 1,500$	NLTK + sklearn
NER Histogram	Named Entity Recognition	9	spaCy
GloVe	Pre-trained embeddings	300	Stanford NLP
Doc2Vec	Trained document vectors	300	Gensim



**Thank You!**